

# Marley Akonnor HW 8

Marley Akonnor

12/3/2021

```
#install.packages("readr")
#install.packages("randomForest")
#install.packages("nnet")
library(readr)
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
library(nnet)

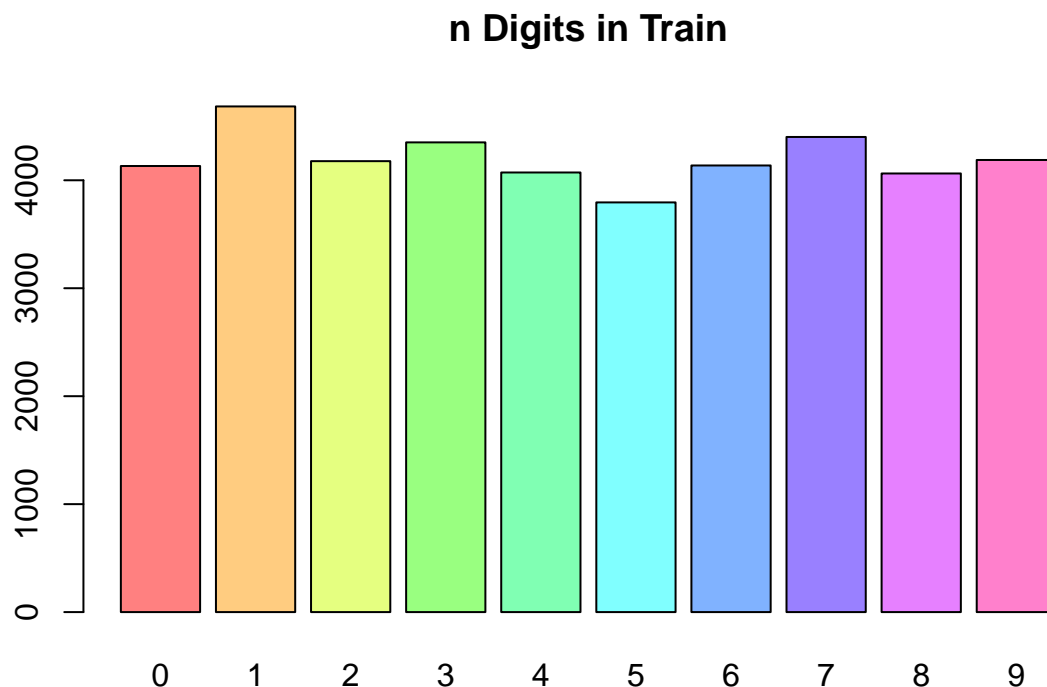
setwd("/Users/m/Documents/M.S Syracuse Data Science/Courses/IST 707 - Data Mining - Machine Learning/Homework")
train_orig <- read_csv("digit_train.csv")

## Rows: 42000 Columns: 785
## -- Column specification -----
## Delimiter: ","
## dbl (785): label, pixel0, pixel1, pixel2, pixel3, pixel4, pixel5, pixel6, pi...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
test_orig <- read_csv("digit_test.csv")

## Rows: 28000 Columns: 784
## -- Column specification -----
## Delimiter: ","
## dbl (784): pixel0, pixel1, pixel2, pixel3, pixel4, pixel5, pixel6, pixel7, p...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
# save the training labels
train_orig_labels <- train_orig[, 1]
train_orig_labels <- as.factor(train_orig_labels$label)
summary(train_orig_labels)

##      0      1      2      3      4      5      6      7      8      9
## 4132 4684 4177 4351 4072 3795 4137 4401 4063 4188

barplot(table(train_orig[,1]), col=rainbow(10, 0.5), main="n Digits in Train")
```



There is around 4000 observations for each digit. Each row has 784

(pixels) which form a 28x28 image. Let's see what the handwritten digits look

like by plotting them. Here is a function to plot a selection of digits from

the train dataset.

```
plotTrain <- function(images, ds, labels){
  op <- par(no.readonly=TRUE)
  x <- ceiling(sqrt(length(images)))
  par(mfrow=c(x, x), mar=c(.1, .1, .1, .1))


  for (i in images){ #reverse and transpose each matrix to rotate images
    m <- matrix(data.matrix(ds[i,-1]), nrow=28, byrow=TRUE)
    m <- apply(m, 2, rev)
    image(t(m), col=grey.colors(255), axes=FALSE)
    text(0.05, 0.2, col="white", cex=1.2, labels[i])
  }
  par(op) #reset the original graphics parameters
}

test_orig_36 <- test_orig[1:36,]
test_orig_36$label <- c(2,0,9,0,3,7,0,3,0,3,5,7,4,0,4,3,3,1,9,0,9,1,1,5,7,4,2,7,4,7,7,5,4,2,6,2)
```


Now let's use this function to look at the first 36 images. You can look at

many images if you wanted too, e.g., `plotTrain(1001:1100)`

```
plotTrain(1:36, train_orig, train_orig$label)
```



```
plotTrain(1:36, test_orig_36, test_orig_36$label)
```



first we are going to try a random forest

```
numTrees <- 25
```

Train on entire training dataset and predict on the test

```
startTime <- proc.time()
rf <- randomForest(train_orig[-1], train_orig_labels, xtest=test_orig, ntree=numTrees)
proc.time() - startTime
```

```
##      user  system elapsed
## 142.996   1.097  144.275
```

user system elapsed


165.52 6.52 207.74

```
print(rf)
```

```
##
## Call:
## randomForest(x = train_orig[-1], y = train_orig_labels, xtest = test_orig,      ntree = numTrees)
##               Type of random forest: classification
##               Number of trees: 25
```

```
## No. of variables tried at each split: 28
##
##      OOB estimate of  error rate: 6.44%
## Confusion matrix:
##      0      1      2      3      4      5      6      7      8      9 class.error
## 0 4030      0      8      7      7      9     31      2     29      9 0.02468538
## 1      0 4585     29     15      9      9      8     10     16      3 0.02113578
## 2     33     19 3902     41     34     20     22     55     34     17 0.06583672
## 3     12     10    101 3953     10    111     14     38     70     32 0.09147322
## 4      5      11     20      6 3812     14     27     16     21    140 0.06385069
## 5     23      9     15    146     22 3441     42      8     55     34 0.09328063
## 6     41      7     18      8     19     53 3962      1     24      4 0.04230118
## 7      3     19     67     18     32      4      3 4141     15     99 0.05907748
## 8     22     33     57     95     36     62     29     13 3646     70 0.10263352
## 9     20     10     23     58     96     43      7     62     47 3822 0.08739255
```

```
plot(rf,type="l")
```



## output predictions for submission

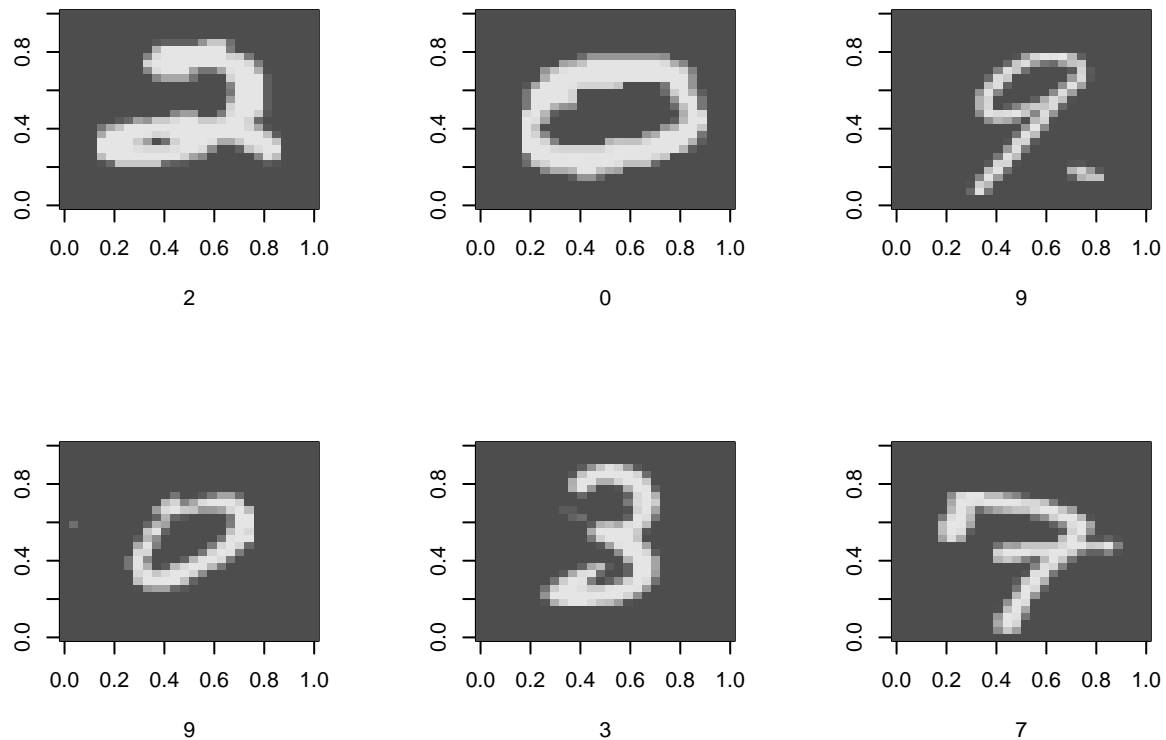
```
predictions <- data.frame(ImageId=1:nrow(test_orig),
  Label=levels(train_orig_labels)[rf$test$predicted])
head(predictions)
```

```
##   ImageId Label
## 1        1     2
## 2        2     0
## 3        3     9
## 4        4     9
```

```
## 5      5      3
## 6      6      7


rotate <- function(x) t(apply(x, 2, rev)) # reverses (rotates the matrix)

par(mfrow=c(2,3)) # Plotting in 2*3 format (random forest)
lapply(1:6,
  function(x) image( #nrow = 28 because this is 28 pixel image
    rotate(matrix(unlist(test_orig[x,]),nrow = 28,byrow = T)),
    col=grey.colors(255),
    xlab=predictions[x,2]
  )
)
```



```
## [[1]]
## NULL
##
## [[2]]
## NULL
##
## [[3]]
## NULL
##
## [[4]]
## NULL
##
## [[5]]
## NULL
##
## [[6]]
## NULL
```

```
plotTrain(1:36, test_orig_36, predictions$Label)
```



```
predictions_36 <- predictions[1:36,2]
accuracy <- mean(predictions_36 == test_orig_36$label)
print(paste('Accuracy:', accuracy))
```

```
## [1] "Accuracy: 0.9444444444444444"
```

split the training data into train and test to do local evaluation

```
set.seed(123)
rows <- sample(1:nrow(train_orig), as.integer(0.7*nrow(train_orig)))
```

Get train and test labels

```
train_labels <- train_orig[rows, 1]
test_labels <- train_orig[-rows, 1]
```

convert the labels to factors

```
train_labels <- as.factor(train_labels$label)
```

## custom normalization function

```
normalize <- function(x) {  
  return(x / 255)  
}
```

## create the train and test datasets and apply normalization

```
train_norm <- as.data.frame(lapply(train_orig[rows, -1], normalize))  
test_norm <- as.data.frame(lapply(train_orig[-rows, -1], normalize))
```

## check a random pixel to see if the normalization worked

```
summary(train_orig$pixel350)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##      0.00   0.00   0.00  89.51  228.00  255.00  
  
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##      0.00   0.00   0.00  89.51  228.00  255.00
```

```
summary(train_norm$pixel350)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## 0.00000 0.00000 0.00000 0.3487 0.8902 1.0000  
  
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## 0.000000 0.000000 0.003922 0.350500 0.890200 1.000000
```

```
summary(test_norm$pixel350)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## 0.00000 0.00000 0.01176 0.35629 0.90196 1.00000
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.  
0.0000 0.0000 0.0000 0.3521 0.9059 1.0000
```

## create the class indicator matrix

```
train_labels_matrix = class.ind(train_labels)  
head(train_labels)
```

```
## [1] 7 6 3 3 1 4  
## Levels: 0 1 2 3 4 5 6 7 8 9
```

```
[1] 2 2 9 2 8 7
```

```
Levels: 0 1 2 3 4 5 6 7 8 9
```

```
head(train_labels_matrix)
```

```
##      0 1 2 3 4 5 6 7 8 9
```



```
## [1,] 0 0 0 0 0 0 0 1 0 0
## [2,] 0 0 0 0 0 0 1 0 0 0
## [3,] 0 0 0 1 0 0 0 0 0 0
## [4,] 0 0 0 1 0 0 0 0 0 0
## [5,] 0 1 0 0 0 0 0 0 0 0
## [6,] 0 0 0 0 1 0 0 0 0 0
```

**0 1 2 3 4 5 6 7 8 9**

```
[1,] 0 0 1 0 0 0 0 0 0 0
[2,] 0 0 1 0 0 0 0 0 0 0
[3,] 0 0 0 0 0 0 0 0 0 1
[4,] 0 0 1 0 0 0 0 0 0 0
[5,] 0 0 0 0 0 0 0 0 1 0
[6,] 0 0 0 0 0 0 0 1 0 0
```

**train model**

```
set.seed(123)
startTime <- proc.time()
nn = nnet(train_norm, train_labels_matrix, size = 1, softmax = TRUE)
```

```
## # weights: 805
## initial value 71562.488017
## iter 10 value 67404.101410
## iter 20 value 64669.742937
## iter 30 value 59259.471988
## iter 40 value 54208.857464
## iter 50 value 52919.684509
## iter 60 value 52486.942418
## iter 70 value 52188.376159
## iter 80 value 52043.260352
## iter 90 value 51942.104544
## iter 100 value 51873.801939
## final value 51873.801939
## stopped after 100 iterations
```

**# weights: 805**

```
initial value 71631.780715
iter 10 value 64946.553191
iter 20 value 57294.824640
iter 30 value 55912.804141
iter 40 value 54648.757612
iter 50 value 53950.781576
iter 60 value 52927.199756
```

```
iter 70 value 52291.634751
iter 80 value 51967.602466
iter 90 value 51774.654787
iter 100 value 51643.951402
final value 51643.951402
stopped after 100 iterations
```

```
proc.time() - startTime
```

```
##    user  system elapsed
## 38.722   0.352  39.115
```

```
user system elapsed
```

```
46.97 0.13 47.54
```

```
nn
```

a 784-1-10 network with 805 weights

options were - softmax modelling

This is just to try out the nnet function. One hidden node is mostly likely

not enough for the model. “softmax” should be set to TRUE when performing

classification. The default maximum number of iterations is 100. The algorithm

did not converge before reaching the maximum. It ran for 46 seconds.

get predictions

```
pred = predict(nn, test_norm, type="class")
cbind(head(pred), head(test_labels))
```

```
##    head(pred) label
## 1           7     9
## 2           1     2
## 3           1     0
## 4           1     6
## 5           7     9
## 6           1     1
```

head(pred) label

1 1 8

2 1 3

3 1 8

4 1 0

5 1 3

6 7 4

The first six predictions do not look very good.







evaluate the model

```
accuracy <- mean(pred == test_labels)
print(paste('Accuracy:', accuracy))
```

```
## [1] "Accuracy: 0.207681930005555"
```

```
## [1] "Accuracy: 0.206174113165622"
```

```
par(mfrow=c(2,3)) # Plotting in 2*3 format (neural net)
lapply(1:6,
  function(x) image( #nrow = 28 because this is 28 pixel image
    rotate(matrix(unlist(test_orig[x,]),nrow = 28,byrow = T)),
    col=grey.colors(255),
    xlab=pred[x]
  )
)
```



```
## [[1]]
## NULL
##
## [[2]]
## NULL
##
## [[3]]
## NULL
##
## [[4]]
## NULL
##
## [[5]]
## NULL
##
## [[6]]
## NULL
```